

## ChatGPT

28-35 minutes

# FFmpeg Best Practices for Installation, Filtering, Encoding, and Usage

## Building and Installing FFmpeg from Source (Ubuntu)

Building FFmpeg from source on Ubuntu ensures you have the latest version with all necessary codec libraries and optimizations. The official guide recommends a **local, non-system installation** so as not to conflict with system packages, using a dedicated build directory (e.g. `~/ffmpeg_build`) [trac.ffmpeg.org](https://trac.ffmpeg.org). Key steps include: installing development tools and dependencies (compiler, autotools, etc.), installing an assembler like YASM or NASM (required for assembly-optimized codecs), and obtaining or compiling codec libraries such as x264, x265, libvpx, libaom, and audio codecs. For example, you should install or build libraries for H.264/H.265 (libx264, libx265), VP8/VP9 (libvpx), AV1 (libaom), AAC (such as libfdk\_aac or use the native encoder), MP3 (libmp3lame), OPUS (libopus), Vorbis (libvorbis), and others needed by your project. When configuring FFmpeg, enable these components so they get compiled in. A typical configure invocation enables GPL and non-free codecs if needed, and links all desired libs, for example: `--enable-gpl --enable-libass --enable-libfreetype --enable-libmp3lame --enable-libopus --enable-libtheora --enable-libvorbis --enable-libvpx --enable-libx264 --enable-libx265 --enable-nonfree` [gist.github.com](https://gist.github.com). This ensures FFmpeg supports subtitles rendering (libass/freetype for text), popular video/audio formats, and codecs needed for H.264/H.265 encoding. It's also common to include `--enable-shared` if you need FFmpeg libraries as `.so`, or use `--enable-static --pkg-config-flags="--static"` for a static build. After configuration, compile with `make -j$(nproc)` to use all CPU cores, and then `make install` to install to the chosen prefix (which can be your home directory to avoid system directories). Add the FFmpeg **bin** directory to your PATH so that your wrapper can invoke the newly built ffmpeg.

**Optimizations:** Compile with optimizations for your CPU. For instance, adding `-extra-cflags="-march=native"` (on x86) or `-mcpu=native` (on ARM) can allow GCC to optimize for the host CPU [amperecomputing.com](https://amperecomputing.com). Ensure YASM/NASM is present so codec libraries use hand-optimized assembly; without it, performance will suffer. If building for a specific use-case, you can disable components you

don't need (to speed up build or reduce size), but generally it's safest to include all commonly used codecs for flexibility. Keep FFmpeg updated to benefit from performance improvements and new filters - using the latest stable release or FFmpeg Git master is a best practice for a wrapper that may need cutting-edge features.

**Common Pitfalls (Installation):** Avoid missing dependencies - if libx264 or others are not found, FFmpeg's configure will silently disable those encoders, leading to *"unknown encoder"* errors later. Always review the output of ./configure to verify desired features are included. For example, if **yasm** is not installed, x264/x265 assembly will be disabled or compilation may fail. When using non-free codecs like libfdk\_aac, remember to configure with `--enable-nonfree`[gist.github.com](https://gist.github.com), otherwise the build will not include them. Also, be mindful of license implications: enabling non-free codecs means the resulting binary is not freely redistributable. Another pitfall is installing your custom build over a distro ffmpeg - it's better to keep them separate (use /usr/local or a private directory) to avoid package conflicts. Finally, if you had a previous FFmpeg in PATH, run `hash -r` or rehash the shell so the new ffmpeg is used[gist.github.com](https://gist.github.com). For a wrapper project, ensure the environment calls the correct FFmpeg binary and not an outdated version.

## Constructing Efficient Filter Chains

When processing video, using FFmpeg's **filter chain** (`filtergraph`) correctly is crucial for efficiency and quality. Filters can be applied in a chain (output of one feeds into the next, separated by commas in a `-vf` or `-filter_complex` description) or in parallel with complex filter graphs (separate chains separated by semicolons with labeled inputs/outputs). The **logic and order** of filters significantly affect the output. A general best practice is to perform **deinterlacing first**, if the source is interlaced, then apply other processing filters, and do **scaling (resizing) last**[superuser.com](https://superuser.com). This is because scaling early can magnify artifacts or noise, and deinterlacing needs to work on the original interlaced fields. As one expert notes, *"start with yadif or any other deinterlacing, then all other filters, and finally end with the scaling filter."*[superuser.com](https://superuser.com). Similarly, if you need to **crop** (e.g. remove black bars) and then scale, do the cropping first, then scale - this preserves maximum detail in the area of interest and avoids wasting effort scaling parts of the frame that will be cut off. **Denoising** filters (such as `hqdn3d`, `nlmeans`, etc.) are often best applied before scaling as well, so that the denoiser works on the original detail (denoising after upscaling would have to process more pixels and could blur the image, and denoising after downscaling might be less effective if noise is scaled down along with the image). In summary, design the filter chain in natural order: *source → decode → deinterlace → crop/pad*

→ *denoise/color adjust* → *scale* → *subtitles/overlay* → *output*. Final overlays like subtitles or watermarks are usually added **after scaling** to ensure they render at the correct final resolution (you generally don't want to scale subtitles or logos along with the video).

Managing **color space and pixel format** is also important in filter chains. Many filters require a certain pixel format (e.g. some work only in YUV, not RGB). It's good practice to explicitly insert a format/color space conversion when needed using the `format` or `colorspace` filters. For example, if you're scaling a high-resolution video and then encoding to H.264 for broad compatibility, you might do `-vf "colorspace=bt709:all=bt601-6-5-4,scale=1280:-1:flags=lanczos,format=yuv420p"` - this ensures the color primaries are converted correctly and the output is in **yuv420p** (the standard 4:2:0 format H.264 expects for widest device support). In simpler cases, you can also just specify `-pix_fmt yuv420p` as an output option, and FFmpeg will insert the conversion automatically. Always preserve the proper color range when filtering (for instance, use `scale` with `in_range/out_range` or the `colormatrix` filter if converting between full and limited range).

For **frame rate adjustments**, FFmpeg offers filters like `fps` (to duplicate/drop frames to a target rate) or more advanced interpolation filters like `minterpolate` (which uses motion estimation to synthesize intermediate frames for smoother motion when converting, say, 30fps to 60fps). If using frame interpolation, be aware it is CPU intensive. It's often performed after spatial filters and at the final output resolution. For example, you might `denoise` and `scale` a video to 720p, then apply `minterpolate` to go from 30→60fps, then encode. Doing interpolation at a lower resolution (after scaling down) is much faster than at full source resolution, and the visual quality benefit of interpolation is usually judged at the playback resolution. Thus, integrate such filters appropriately in the chain considering performance trade-offs.

To construct filter chains efficiently: use **simple filtergraphs** (with `-vf` for one input stream) when possible, chaining multiple filters with commas. This ensures the video is decoded and passed through all filters in one go. For example: `-vf "deinterlace,scale=1280:-1,format=yuv420p"` applies three filters sequentially on the input. If you accidentally use multiple `-vf` options on the same output, only the last one takes effect (each `-vf` overrides the previous for that output), which is a common mistake. For more complex scenarios (like filtering multiple inputs or splitting and recombining streams), use `-filter_complex` with labeled pads. Ensure filtergraph syntax is correct and **escape any special characters** in shell contexts (e.g., use quotes around the filter description, and escape characters like `,` or `:` if needed in a shell script). Tools like *alfg's FFmpeg Commander* can help visually build filter graphs, but under the hood they follow the same principles.

**Common Pitfalls (Filter Chains):** One pitfall is applying filters in the wrong order. For instance, **upsampling too early** will enlarge noise and make subsequent denoising or compression less effective [superuser.com](https://superuser.com). Always try to perform cleanup (deinterlace, denoise) at the native resolution and scale at the end. Another common mistake is using separate filter options for multiple filters (e.g. `-vf scale=1280:-1 -vf transpose=1`); only the last one will run - instead, combine them (`-vf "transpose=1,scale=1280:-1"`). Forgetting to enforce pixel format can lead to compatibility issues: for example, if you don't specify `yuv420p`, FFmpeg might output 4:4:4 or 4:2:2 colors which some players can't handle, or 10-bit depth if your pipeline inadvertently stays in high bit-depth - always downconvert to 8-bit 4:2:0 for standard outputs unless you explicitly want high bit-depth. Also be careful with **escaping and quoting**: complex filtergraphs containing symbols `[]`, `;`, etc., must be quoted or escaped in most shells. If your wrapper builds command strings, ensure it quotes filter arguments properly to avoid "Unable to parse filter graph" errors. Lastly, check each filter's documentation for default behaviors - e.g., the `scale` filter defaults to bicubic scaling; you might choose `-vf scale=1280:-1:flags=lanczos` for sharper resizing if desired. Testing your filter chain on a short sample video can verify the order and result before processing large files.

## Optimal H.264 Encoding Practices

H.264 (AVC) is one of the most widely used codecs, and **libx264** is the FFmpeg encoder for it. To get the best results, use the proper encoding parameters rather than just default bitrate. The recommended approach is to use **constant quality** mode with the `-crf` (Constant Rate Factor) parameter. CRF ranges from 0 (lossless) to 51 (worst quality). Lower CRF means higher quality (and larger filesize). The default is 23, and *visually lossless* quality is typically around **CRF 18** [gist.github.com](https://gist.github.com). In practice, a CRF between about 18-28 is a sensible range for most uses [gist.github.com](https://gist.github.com) - use 18-20 for very high quality, ~23 for standard quality, and higher values for more compression when some quality loss is acceptable. Unlike a fixed bitrate, CRF will let the bitrate vary based on content complexity to maintain consistent quality. This is usually preferred for offline encoding because it frees you from guessing an optimal bitrate. If you do need a specific output size or bitrate (e.g., for streaming or strict bandwidth limits), use **two-pass encoding** with a target bitrate. In two-pass mode, you run `ffmpeg -pass 1 -c:v libx264 -b:v [target] -preset [preset] -an -f mp4 /dev/null` (no audio, discard output) then `-pass 2` with the real output. This ensures the encoder can allocate bits optimally. However, for most use cases, CRF with a reasonable value is simpler and delivers great results.

The **preset** is another crucial setting. x264's preset tunes the encoding speed vs compression efficiency. They range from ultrafast, superfast, veryfast, faster, fast, **medium** (default), slow, slower, **veryslow**, to placebo. Slower presets give better compression (quality/bitrate) but take more time. For example, using -preset slower or veryslow can shave some file size off or improve quality at the same bitrate compared to medium [gist.github.com](https://gist.github.com). A common practice is to use the slowest preset that you can afford time-wise for your application. For a one-time high quality encode, veryslow or slower is great. For real-time or quick turnaround, you might need fast or veryfast. **Do not mix up preset with quality** - preset doesn't affect visual quality directly, it affects how efficiently the encoder finds compression. So you use preset in conjunction with CRF: e.g. -c:v libx264 -crf 22 -preset slow. Changing preset from medium to slow might let you drop CRF by a point for the same output size, but there are diminishing returns (e.g., placebo is usually not worth the huge increase in encoding time for a negligible improvement).

When targeting specific **devices or compatibility** requirements, set the **profile and level**. H.264 has profiles like Baseline, Main, High, etc. Baseline Profile lacks features like B-frames and is used for older or low-power devices. Main and High offer better compression. Most modern devices (PCs, smartphones, etc.) support High Profile, so generally use -profile:v high for best quality. However, if you need broad compatibility with very old devices (early smartphones, older embedded systems), you might choose -profile:v baseline or main. For example, **Baseline@L3.0** is often needed for older Android phones or hardware decoders [gist.github.com](https://gist.github.com). As the FFmpeg Wiki notes, you can set profiles with -profile:v baseline or -profile:v main, but "*most modern devices support the more advanced High profile.*" [trac.ffmpeg.org](https://trac.ffmpeg.org). The **level** (e.g. 3.0, 4.0, 4.1) defines constraints like max bitrate and resolution; choose a level that accommodates your video resolution/frame-rate (e.g. 4.0 for 1080p30, 4.2 for 1080p60, 5.1 for 4K, etc.) and that your target devices support. If unsure, you can omit level and let x264 automatically set the minimum level needed (it prints the level in the output log). Always combine profile/level settings with ensuring the correct pixel format. For H.264 in MP4, that means **YUV 4:2:0 8-bit**: use -pix\_fmt yuv420p unless you have a specific reason to keep a higher format. This ensures compatibility with all players (for instance, if you encode to YUV444 or even YUV420 10-bit, many decoders - including default QuickTime on macOS or older mobile devices - won't play it). By default, libx264 will output yuv420p for an 8-bit build, but if your input is RGB or higher bit depth, explicitly forcing yuv420p is wise. The common "**compatibility recipe**" for web video is: -c:v libx264 -crf 20 -preset slow -pix\_fmt yuv420p -profile:v baseline -level 3.0 -movflags +faststart to cover all bases [gist.github.com](https://gist.github.com) [gist.github.com](https://gist.github.com).

Speaking of **-movflags +faststart**: this is a muxer option for MP4/MOV outputs that moves the file's index (moov atom) to the beginning of the file. It is recommended for any video that will be streamed or played progressively (e.g. in a browser)[gist.github.com](https://gist.github.com). Without it, the video might not start playing until fully downloaded. It doesn't affect the encoding itself, but it optimizes the MP4 for web delivery.

Other H.264 encoding tips: use **-tune** options if appropriate - for example, **-tune film** (if your video has film grain and you want to preserve it), **-tune animation** (for cartoons), or **-tune grain**. These adjust encoder decisions to better suit those content types. Do note that using **-tune** is optional; only use one if it matches your content because it will trade off compression efficiency in other areas (e.g., **-tune film** disables some psy optimizations to avoid adding artifacts to grain, but that can increase bitrate). If you want **lossless** H.264, you can use **-crf 0** or **-qp 0** (they are equivalent for x264; CRF 0 means lossless). This will result in huge files, so it's rarely used except for intermediary files. A more practical near-lossless is CRF 18 as mentioned (almost indistinguishable from source in most cases[gist.github.com](https://gist.github.com)). Also, remember to set **-c:a** for audio (e.g., **-c:a aac -b:a 192k** for AAC audio at 192 kbps) or use **-c:a copy** to copy original audio if it's already a good format.

**Common Pitfalls (H.264 Encoding):** A common mistake is **bitrate overspecification** - for example, using **-crf** and also **-b:v**. If you specify both, x264 will actually use CRF quality but with a bitrate cap (constrained quality). This is advanced; if you don't intend that, avoid mixing CRF and bitrate in one command. Decide on one rate-control method. Another pitfall is forgetting to set **-pix\_fmt yuv420p** for compatibility - you might end up with a perfectly encoded video that won't play on your iPhone due to it being yuv444p. Always downscale chroma for final output unless you know the target supports it. Using **too slow a preset** can be a problem in a different sense: e.g., **-preset placebo** can take an extraordinarily long time for negligible improvement over **veryslow**. Professional workflows rarely go beyond **veryslow**. On the other end, using **ultrafast** preset with a low CRF can degrade quality because **ultrafast** disables many compression tools; you might get larger files that still look worse. It's about balance. Make sure the **profile** and **level** are set appropriately: if you force **-profile high** but also **-level 3.0**, you might inadvertently make the video unplayable on a device that only supports **Baseline@3.0** (because High profile tools aren't supported even if level is low). In general, if targeting old devices, use **Baseline**; otherwise leave profile as default (which will choose High automatically) for best efficiency, and only limit level if needed. Lastly, **not using faststart** for MP4 when needed is a minor pitfall - your file will still play, but users might experience buffering; so it's a good practice to always include **+faststart** for MP4 outputs meant for streaming. For a wrapper project, it's wise to expose

CRF, preset, profile, etc., as configurable parameters with safe defaults (e.g., default to crf 23, preset medium, high profile) and document these for the end users.

## FFmpeg Command Usage and Key Concepts

Using FFmpeg effectively in a wrapper requires understanding how its command-line syntax and options work. The **ffmpeg command structure** is: `ffmpeg [global_options] -i input1 [input_options] -i input2 ... [output_options] output1 [output_options] output2 ....` Order matters greatly in FFmpeg's command line. Options apply to the next relevant file or stream. In general, **global options** (applying to the whole process, like `-y` to overwrite, or logging level) should come first. Then you list all inputs with their options, then outputs with their options. The documentation emphasizes: *"Do not mix input and output files - first specify all input files, then all output files. Also do not mix options which belong to different files. All options apply **only** to the next input or output file and are reset between files."* [ffmpeg.org](http://ffmpeg.org). For example, if you want to set an option for an output, put it *before* that output's filename, not after. If you want to apply an input option (like `-r` for raw input frame rate or `-ss` to seek), it must come right before the `-i` for that input. A common error is to put an output option too early and FFmpeg interprets it as an input option or vice versa. Following the correct order avoids unexpected behaviors.

**Stream selection and mapping:** By default, FFmpeg will automatically pick one stream of each type (video, audio, subtitles) from the inputs to include in each output. The selection defaults are documented: it picks the *best* quality streams (highest resolution video, highest channel count audio) [ffmpeg.org](http://ffmpeg.org) if you don't explicitly choose. If you have complex inputs or want specific streams, use the `-map` option. FFmpeg *"provides the `-map` option for manual control of stream selection in each output file"*, otherwise it performs automatic selection [ffmpeg.org](http://ffmpeg.org). For instance, `-map 0:v:0 -map 1:a:0` would take the first video stream from input 0 and the first audio from input 1. It's good practice in a multi-input or multi-output scenario to use `-map` rather than relying on the default, to avoid surprises (like the wrong audio language being picked). If you want to exclude a type of stream, you can use `-vn` (no video), `-an` (no audio), `-sn` (no subtitles) on the output. These can be combined with mapping. Note that if you use complex filtergraphs (`filter_complex`), those implicitly create new streams that you'll need to map manually (or they get added to the first output by default which can be confusing). In a wrapper, if the user expects to, say, drop the audio, you should include `-an` or skip mapping audio streams accordingly, instead of encoding a silent audio track.



**Codec and format selection:** Explicitly set codecs for outputs using `-c:v`, `-c:a`, etc. If you don't, FFmpeg will try to pick a default encoder based on the output file format. For example, if output is `.mp4` and you didn't specify `-c:v`, historically FFmpeg might choose the `mpeg4` codec (not to be confused with H.264) or another default, which may not be what you want. It's safer to always specify the codec (like `-c:v libx264 -c:a aac` for MP4). Use `-c copy` (stream copy) when you want to mux an input stream to output without re-encoding - this is great for copying audio or video that is already in the desired format. For instance, if your input has AAC audio and you are making an MP4, `-c:a copy` will just transfer the AAC stream. Be mindful that when stream copying, the format must support that codec or it will error (e.g., you can't copy a DTS audio stream into an MP4 because MP4 doesn't support DTS). That relates to **multiplexing**: the container format (specified by output file extension or `-f`) must be able to hold the chosen codecs. FFmpeg usually knows the common ones, but if you attempt something unsupported, it will throw an error. A best practice is to use common, widely compatible codecs for each format (H.264/AAC for MP4, etc.) unless you have a specific reason to do otherwise.

**Pixel formats and color details:** As discussed, using `-pix_fmt` for output is important if you need a specific pixel format. You can list supported pixel formats with `ffmpeg -pix_fmts`. If a chosen encoder cannot handle the input pixel format, FFmpeg will try to insert a converter. Sometimes it fails if there's ambiguity, so it's better to request it. For example, for GIF output, you'd use `-pix_fmt pal8`. For `libx264`, stick to `yuv420p` (8-bit) unless you compiled `x264` for 10-bit (then the encoder expects `yuv420p10le`, and you'd have to provide that). Ensuring the correct colorspace (e.g., using the `-colorspace`, `-color_primaries`, and `-color_trc` metadata options or filters) is also a consideration for high-quality workflows, but for standard uses FFmpeg will carry over or set reasonable defaults for these based on input.

**Metadata handling:** FFmpeg can read and write metadata tags (like title, artist, encoding software, etc.) as well as stream metadata (like language for audio tracks). By default, when transcoding, FFmpeg will not automatically copy all metadata from input to output (it does copy some like language, but not all tags). If you need to preserve metadata, you have a couple options: use `-map_metadata 0` to copy global metadata from input 0 to output, and similarly `-map_metadata:s:a:0 0:s:a:0` for stream metadata, etc. You can also set specific metadata with `-metadata` options. For example: `-metadata title="My Video"` sets a title tag [ffmpeg.org](https://ffmpeg.org), or `-metadata:s:a:0 language=eng` to tag the first audio stream as English [ffmpeg.org](https://ffmpeg.org). Your wrapper could expose fields to set these. Also, if dealing with chapters, the `-map_chapters 0` option would copy chapters from the first input. It's important to note that when you stream copy a track, its metadata (like language, codec name) is usually preserved; when you transcode, you may need to re-set some metadata. Make sure to **delete or update**



**any metadata** that no longer applies (e.g., the encoder tag will say Lavf/FFmpeg by default – you can override `-metadata encoding_tool="MyApp 1.0"` if desired).

**Hardware acceleration:** FFmpeg supports various hardware accelerations for both decoding and encoding. For decoding, you can use `-hwaccel` to enable a hardware decoder for the input. For instance, `-hwaccel cuda -hwaccel_output_format cuda` can be used with NVIDIA GPUs to speed up decoding of compatible codecs. This can dramatically reduce CPU usage for high-res videos. The ffmpeg documentation states `-hwaccel` will “*Use hardware acceleration to decode the matching stream(s).*”[ffmpeg.org](https://ffmpeg.org). Common values are `cuda`, `dxva2` (Windows), `vaapi` (Linux Intel/AMD), etc. In a wrapper, you might let advanced users toggle hardware decode if available. For encoding, you have to choose a hardware-supported encoder, e.g. `-c:v h264_nvenc` for NVIDIA NVENC, `-c:v hevc_qsv` for Intel QuickSync HEVC, `-c:v h264_amf` for AMD AMF, etc. These encoders run on the GPU and can be much faster, at the cost of some quality (generally, x264 slow preset still beats NVENC in quality per bitrate, but NVENC is real-time and very convenient). Ensure your FFmpeg build has these enabled (e.g., configure with `--enable-nvenc`, `--enable-libmfx` for QSV, etc., and that the system has the proper drivers). If using hardware encoders, also consider their **options**: for example, NVENC has its own preset system (gpu-specific) and rate control modes. Also note that some filters can run on hardware (like `scale_npp` on NVIDIA, or `deinterlace_qsv` on Intel). If your pipeline is heavy on CPU filters, the benefit of hardware decode/encode might diminish due to CPU-GPU memory transfers. A balanced approach is needed.

**Multiplexing and output:** Once encoding is done, FFmpeg’s muxer writes the output file. To ensure a clean output, always use an appropriate container file extension or explicitly `-f format`. For instance, to output an MPEG-TS stream you’d do `-f mpegts output.ts`. For MP4, just naming file `.mp4` is enough. If producing fragmented MP4 or DASH, there are specific flags (beyond scope here). One important tip: for **live streaming**, some formats require specific muxer settings (e.g., using `-f flv` for RTMP or `-f mpegts` for UDP multicast, etc.). These might not apply to a wrapper intended for file transcoding, but it’s good to be aware. When creating outputs with multiple streams (say two audio tracks or an additional subtitle), ensure you map and encode each, and the output container can hold them (MKV is very flexible with streams, MP4 has some limitations like at most one video track generally and certain subtitle formats not supported). If including subtitles, you may need to choose between embedding as SRT text vs. burning into video via filter (burning in makes them part of the video, which is necessary if the target doesn’t support subtitle streams).

**Error handling:** Robust error handling is vital in a wrapper. FFmpeg will return a non-zero exit code on failure. Make sure to capture `stderr` output; FFmpeg

prints progress and errors to stderr by default. Monitor for messages like "Error" or "[warning]" in the output. For example, if a filter name is wrong, FFmpeg might error out with "No such filter" or "Error initializing filters". If an encoder is not found (not compiled in), you'll see "*Unknown encoder 'xyz'*". Your wrapper should detect these and report them clearly (possibly suggesting to check build configuration if a codec is missing). To aid debugging, you can run ffmpeg with higher verbosity (-loglevel debug or -report to create a log file). In automated use, it's common to add -hide\_banner to suppress the library version banner for cleanliness, but keep the warnings and errors. This makes parsing easier. Additionally, consider handling **interruptions** - if the user cancels, you might send a SIGINT to FFmpeg (which will attempt a safe shutdown of muxer). If encoding multiple files in a batch, handle failures on one file by continuing to the next appropriately.

**Common Pitfalls (Usage/Commands):** Many pitfalls here revolve around option ordering and selection. One classic mistake is placing an output option too early, e.g., `ffmpeg -c:v libx264 -i input.mp4 output.mp4`. In this command, `-c:v libx264` is actually interpreted as an **input** option (before `-i`), which does nothing, and thus the output might use a default codec (MPEG-4) instead of x264. The correct order is `ffmpeg -i input.mp4 -c:v libx264 output.mp4`. Always remember the rule: options apply to the next file. Another pitfall is not using `-map` in complex scenarios: if you have multiple audio tracks and you transcode, by default FFmpeg might only pick one. People often wonder "where did my second audio go?" - it wasn't mapped, so it wasn't included. Your wrapper could automatically map all streams if the goal is to preserve them, or provide user control. Also, **forgetting to override defaults** can bite you: for instance, if you remux to MP4 without `-c:a copy`, and the input had MP3 audio, FFmpeg might choose to re-encode that to AAC by default (because MP3 is not standard in MP4). That might be fine or might be unwanted. It's better to explicitly copy streams when you intend to. Another common error scenario is **mismatched arguments** - e.g. giving an audio encoder option to a video stream or a filter that doesn't exist. FFmpeg will usually error out in those cases. Pay attention to its error messages; they're usually quite descriptive. For example, "*Tag xyz incompatible with output format*" means the codec isn't supported in that container. Or "*At least one output file must be specified*" means you forgot to put the output file argument. In a scripting context, always double-check that file paths are correctly passed (wrap paths in quotes if they may contain spaces, as the script from the FFmpeg guide suggests [gist.github.com](https://gist.github.com)).

In summary, to build a robust FFmpeg wrapper: use the **proper build of FFmpeg** with all needed codecs, construct filter chains in a logical order for quality, use proven encoding settings for H.264 (or any codec in use), and assemble the FFmpeg command with correct option ordering and explicit mappings. Handle

errors gracefully by checking FFmpeg's output and exit code, and provide informative feedback to the user. By following these best practices and avoiding the common pitfalls noted, your FFmpeg wrapper will produce reliable, high-quality results and accommodate a wide range of media processing tasks.

[ffmpeg.org](https://ffmpeg.org)[ffmpeg.org](https://ffmpeg.org)